

## MLP on MNIST Data Set

In [0]:

```
# if you keras is not using tensorflow as backend set "KERAS_BACKEND=tensorflow" use this command
from keras.utils import np_utils
from keras.datasets import mnist
import seaborn as sns
from keras.initializers import RandomNormal
```

In [0]:

```
%matplotlib inline

import matplotlib.pyplot as plt
import numpy as np
import time
plt.style.use('classic')
# https://gist.github.com/greydanus/f6eee59eaf1d90fcb3b534a25362cea4
# https://stackoverflow.com/a/14434334
# this function is used to update the plots for each epoch and error
def plt_dynamic(x, vy, ty, ax, colors=['b']):
    ax.plot(x, vy, 'b', label="Validation Loss")
    ax.plot(x, ty, 'r', label="Train Loss")
    plt.legend()
    plt.grid()
    fig.canvas.draw()
```

## Loading Data

In [9]:

```
# the data, shuffled and split between train and test sets
(X_train, y_train), (X_test, y_test) = mnist.load_data()
```

Downloading data from <https://s3.amazonaws.com/img-datasets/mnist.npz>  
11493376/11490434 [=====] - 1s 0us/step

In [10]:

```
print("Number of training examples :", X_train.shape[0], "and each image is of shape (%d, %d)"%(X_train.shape[1], X_train.shape[2]))
print("Number of training examples :", X_test.shape[0], "and each image is of shape (%d, %d)"%(X_test.shape[1], X_test.shape[2]))
```

Number of training examples : 60000 and each image is of shape (28, 28)  
Number of training examples : 10000 and each image is of shape (28, 28)

In [0]:

```
# if you observe the input shape its 2 dimensional vector
# for each image we have a (28*28) vector
# we will convert the (28*28) vector into single dimensional vector of 1 * 784

X_train = X_train.reshape(X_train.shape[0], X_train.shape[1]*X_train.shape[2])
X_test = X_test.reshape(X_test.shape[0], X_test.shape[1]*X_test.shape[2])
```

In [12]:

```
# after converting the input images from 3d to 2d vectors

print("Number of training examples :", X_train.shape[0], "and each image is of shape (%d)"%(X_train.shape[1]))
print("Number of training examples :", X_test.shape[0], "and each image is of shape (%d)"%(X_test.shape[1]))
```

Number of training examples : 60000 and each image is of shape (784)  
Number of training examples : 10000 and each image is of shape (784)

In [13]:

```
# An example data point
print(X_train[0])
```

```
[ 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  3  18  18  18 126 136 175  26 166 255
247 127  0  0  0  0  0  0  0  0  0  0  0  0  0  30  36  94 154
170 253 253 253 253 253 225 172 253 242 195  64  0  0  0  0  0  0
  0  0  0  0  0  49 238 253 253 253 253 253 253 253 251  93  82
 82  56  39  0  0  0  0  0  0  0  0  0  0  0  0 18 219 253
253 253 253 253 198 182 247 241  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  80 156 107 253 253 205  11  0  43 154
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0 14  1 154 253  90  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0 139 253 190  2  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0 11 190 253  70  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  35 241
225 160 108  1  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  81 240 253 253 119  25  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  45 186 253 253 150  27  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0 16  93 252 253 187
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0 249 253 249  64  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  46 130 183 253
253 207  2  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  39 148 229 253 253 253 250 182  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  24 114 221 253 253 253
253 201  78  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  23  66 213 253 253 253 253 198  81  2  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  18 171 219 253 253 253 195
 80  9  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
 55 172 226 253 253 253 244 133 11  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  136 253 253 253 212 135 132 16
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0]
```

In [0]:

```
# if we observe the above matrix each cell is having a value between 0-255
# before we move to apply machine learning algorithms lets try to normalize the data
#  $X \Rightarrow (X - X_{min}) / (X_{max} - X_{min}) = X / 255$ 

X_train = X_train/255
X_test = X_test/255
```

In [15]:

```
# example data point after normlizing
print(X_train[0])
```

```
[0.  0.  0.  0.  0.  0.
 0.  0.  0.  0.  0.  0.
 0.  0.  0.  0.  0.  0.
 0.  0.  0.  0.  0.  0.]
```

[illegible]

[illegible]

In [16]:

```
# here we are having a class number for each image
print("Class label of first image :", y_train[0])

# lets convert this into a 10 dimensional vector
# ex: consider an image is 5 convert it into 5 => [0, 0, 0, 0, 0, 1, 0, 0, 0, 0]
# this conversion needed for MLPs

Y_train = np_utils.to_categorical(y_train, 10)
Y_test = np_utils.to_categorical(y_test, 10)

print("After converting the output into a vector : ", Y_train[0])
```

```
Class label of first image : 5
After converting the output into a vector : [0. 0. 0. 0. 0. 1. 0. 0. 0. 0.]
```

In [0]:

```
# https://keras.io/getting-started/sequential-model-guide/

# The Sequential model is a linear stack of layers.
# you can create a Sequential model by passing a list of layer instances to the constructor:
```

```

# model = Sequential([
#     Dense(32, input_shape=(784,)),
#     Activation('relu'),
#     Dense(10),
#     Activation('softmax'),
# ])

# You can also simply add layers via the .add() method:

# model = Sequential()
# model.add(Dense(32, input_dim=784))
# model.add(Activation('relu'))

###

# https://keras.io/layers/core/

# keras.layers.Dense(units, activation=None, use_bias=True, kernel_initializer='glorot_uniform',
# bias_initializer='zeros', kernel_regularizer=None, bias_regularizer=None,
# activity_regularizer=None,
# kernel_constraint=None, bias_constraint=None)

# Dense implements the operation: output = activation(dot(input, kernel) + bias) where
# activation is the element-wise activation function passed as the activation argument,
# kernel is a weights matrix created by the layer, and
# bias is a bias vector created by the layer (only applicable if use_bias is True).

# output = activation(dot(input, kernel) + bias) => y = activation(WT. X + b)

####

# https://keras.io/activations/

# Activations can either be used through an Activation layer, or through the activation argument s
upported by all forward layers:

# from keras.layers import Activation, Dense

# model.add(Dense(64))
# model.add(Activation('tanh'))

# This is equivalent to:
# model.add(Dense(64, activation='tanh'))

# there are many activation functions ar available ex: tanh, relu, softmax

from keras.models import Sequential
from keras.layers import Dense, Activation

```

In [0]:

```

# some model parameters

output_dim = 10
input_dim = X_train.shape[1]

batch_size = 128
nb_epoch = 20

```

## MLP + RELU activation + adam optimizer + 2 hidden layer

In [14]:

```

# Multilayer perceptron

model_relu = Sequential()
model_relu.add(Dense(392, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNorm
al(mean=0.0, stddev=0.062, seed=None)))
model_relu.add(Dense(196, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.062,
seed=None)))
model_relu.add(Dense(output_dim, activation='softmax'))

```

```

model_relu.summary()

model_relu.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_relu.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))

```

| Layer (type)              | Output Shape | Param # |
|---------------------------|--------------|---------|
| dense_4 (Dense)           | (None, 392)  | 307720  |
| dense_5 (Dense)           | (None, 196)  | 77028   |
| dense_6 (Dense)           | (None, 10)   | 1970    |
| Total params: 386,718     |              |         |
| Trainable params: 386,718 |              |         |
| Non-trainable params: 0   |              |         |

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/tensorflow/python/ops/math\_ops.py:3066: to\_int32 (from tensorflow.python.ops.math\_ops) is deprecated and will be removed in a future version.

Instructions for updating:

Use tf.cast instead.

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 6s 96us/step - loss: 0.2439 - acc: 0.9285 - val\_loss: 0.1123 - val\_acc: 0.9662

Epoch 2/20

60000/60000 [=====] - 5s 90us/step - loss: 0.0896 - acc: 0.9724 - val\_loss: 0.0830 - val\_acc: 0.9744

Epoch 3/20

60000/60000 [=====] - 5s 89us/step - loss: 0.0560 - acc: 0.9825 - val\_loss: 0.0897 - val\_acc: 0.9717

Epoch 4/20

60000/60000 [=====] - 5s 91us/step - loss: 0.0391 - acc: 0.9876 - val\_loss: 0.0745 - val\_acc: 0.9776

Epoch 5/20

60000/60000 [=====] - 5s 90us/step - loss: 0.0279 - acc: 0.9909 - val\_loss: 0.0783 - val\_acc: 0.9793

Epoch 6/20

60000/60000 [=====] - 6s 92us/step - loss: 0.0229 - acc: 0.9930 - val\_loss: 0.0777 - val\_acc: 0.9779

Epoch 7/20

60000/60000 [=====] - 5s 90us/step - loss: 0.0164 - acc: 0.9947 - val\_loss: 0.0774 - val\_acc: 0.9795

Epoch 8/20

60000/60000 [=====] - 5s 90us/step - loss: 0.0172 - acc: 0.9944 - val\_loss: 0.0766 - val\_acc: 0.9793

Epoch 9/20

60000/60000 [=====] - 5s 89us/step - loss: 0.0140 - acc: 0.9953 - val\_loss: 0.0879 - val\_acc: 0.9785

Epoch 10/20

60000/60000 [=====] - 5s 90us/step - loss: 0.0141 - acc: 0.9954 - val\_loss: 0.0784 - val\_acc: 0.9811

Epoch 11/20

60000/60000 [=====] - 5s 90us/step - loss: 0.0111 - acc: 0.9960 - val\_loss: 0.1045 - val\_acc: 0.9755

Epoch 12/20

60000/60000 [=====] - 5s 91us/step - loss: 0.0124 - acc: 0.9961 - val\_loss: 0.1032 - val\_acc: 0.9768

Epoch 13/20

60000/60000 [=====] - 5s 90us/step - loss: 0.0068 - acc: 0.9978 - val\_loss: 0.0961 - val\_acc: 0.9800

Epoch 14/20

60000/60000 [=====] - 5s 90us/step - loss: 0.0082 - acc: 0.9973 - val\_loss: 0.0886 - val\_acc: 0.9820

Epoch 15/20

60000/60000 [=====] - 5s 90us/step - loss: 0.0138 - acc: 0.9951 - val\_loss: 0.0882 - val\_acc: 0.9798

Epoch 16/20

60000/60000 [=====] - 5s 90us/step - loss: 0.0078 - acc: 0.9974 - val\_loss: 0.0876 - val\_acc: 0.9810

Epoch 17/20

60000/60000 [=====] - 5s 90us/step - loss: 0.0058 - acc: 0.9979 -

```

val_loss: 0.0883 - val_acc: 0.9816
Epoch 18/20
60000/60000 [=====] - 5s 90us/step - loss: 0.0044 - acc: 0.9986 -
val_loss: 0.1004 - val_acc: 0.9806
Epoch 19/20
60000/60000 [=====] - 5s 88us/step - loss: 0.0105 - acc: 0.9965 -
val_loss: 0.0868 - val_acc: 0.9828
Epoch 20/20
60000/60000 [=====] - 5s 88us/step - loss: 0.0057 - acc: 0.9982 -
val_loss: 0.0958 - val_acc: 0.9818

```

In [15]:

```

score = model_relu.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, va
lvalidation_data=(X_test, Y_test))


# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.09581032517525569  
Test accuracy: 0.9818



In [16]:


```
w_after = model_relu.get_weights()
```

```
h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)
```

```
fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



## MLP + RELU activation + adam optimizer + 2 hidden layer + Batch normalization and Drop out

In [4]:

```
from keras.layers.normalization import BatchNormalization
from keras.layers import Dropout
```

Using TensorFlow backend.

In [18]:

```
# Multilayer perceptron

model_relu = Sequential()

model_relu.add(Dense(392, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.062, seed=None)))
model_relu.add(BatchNormalization())
model_relu.add(Dropout(0.5))

model_relu.add(Dense(196, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.062,
```



```

seed=None)))
model_relu.add(BatchNormalization())
model_relu.add(Dropout(0.5))

model_relu.add(Dense(output_dim, activation='softmax'))

model_relu.summary()

model_relu.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_relu.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))

```

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow\_backend.py:3445: calling dropout (from tensorflow.python.ops.nn\_ops) with keep\_prob is deprecated and will be removed in a future version.

Instructions for updating:

Please use `rate` instead of `keep\_prob`. Rate should be set to `rate = 1 - keep\_prob`.

| Layer (type)                                | Output Shape | Param # |
|---|--------------|---------|
| dense_7 (Dense)                             | (None, 392)  | 307720  |
| batch_normalization_1 (Batch Normalization) | (None, 392)  | 1568    |
| dropout_1 (Dropout)                         | (None, 392)  | 0       |
| dense_8 (Dense)                             | (None, 196)  | 77028   |
| batch_normalization_2 (Batch Normalization) | (None, 196)  | 784     |
| dropout_2 (Dropout)                         | (None, 196)  | 0       |
| dense_9 (Dense)                             | (None, 10)   | 1970    |

Total params: 389,070  
 Trainable params: 387,894  
 Non-trainable params: 1,176

Train on 60000 samples, validate on 10000 samples

```

Epoch 1/20
60000/60000 [=====] - 8s 136us/step - loss: 0.4503 - acc: 0.8642 -
val_loss: 0.1457 - val_acc: 0.9537
Epoch 2/20
60000/60000 [=====] - 7s 123us/step - loss: 0.2161 - acc: 0.9353 -
val_loss: 0.1105 - val_acc: 0.9658
Epoch 3/20
60000/60000 [=====] - 7s 121us/step - loss: 0.1633 - acc: 0.9502 -
val_loss: 0.0950 - val_acc: 0.9702
Epoch 4/20
60000/60000 [=====] - 7s 121us/step - loss: 0.1396 - acc: 0.9575 -
val_loss: 0.0854 - val_acc: 0.9727
Epoch 5/20
60000/60000 [=====] - 7s 121us/step - loss: 0.1244 - acc: 0.9611 -
val_loss: 0.0790 - val_acc: 0.9759
Epoch 6/20
60000/60000 [=====] - 7s 120us/step - loss: 0.1122 - acc: 0.9652 -
val_loss: 0.0701 - val_acc: 0.9778
Epoch 7/20
60000/60000 [=====] - 7s 121us/step - loss: 0.1055 - acc: 0.9670 -
val_loss: 0.0677 - val_acc: 0.9798
Epoch 8/20
60000/60000 [=====] - 7s 119us/step - loss: 0.0961 - acc: 0.9706 -
val_loss: 0.0689 - val_acc: 0.9795
Epoch 9/20
60000/60000 [=====] - 7s 120us/step - loss: 0.0877 - acc: 0.9729 -
val_loss: 0.0665 - val_acc: 0.9791
Epoch 10/20
60000/60000 [=====] - 7s 119us/step - loss: 0.0830 - acc: 0.9739 -
val_loss: 0.0629 - val_acc: 0.9815
Epoch 11/20
60000/60000 [=====] - 7s 120us/step - loss: 0.0811 - acc: 0.9745 -
val_loss: 0.0637 - val_acc: 0.9806
Epoch 12/20
60000/60000 [=====] - 7s 122us/step - loss: 0.0730 - acc: 0.9773 -
val_loss: 0.0592 - val_acc: 0.9806

```

```

val_loss: 0.0599 - val_acc: 0.9814
Epoch 13/20
60000/60000 [=====] - 8s 127us/step - loss: 0.0738 - acc: 0.9766 -
val_loss: 0.0599 - val_acc: 0.9814
Epoch 14/20
60000/60000 [=====] - 7s 124us/step - loss: 0.0704 - acc: 0.9777 -
val_loss: 0.0604 - val_acc: 0.9813
Epoch 15/20
60000/60000 [=====] - 8s 127us/step - loss: 0.0644 - acc: 0.9793 -
val_loss: 0.0558 - val_acc: 0.9829
Epoch 16/20
60000/60000 [=====] - 8s 126us/step - loss: 0.0646 - acc: 0.9789 -
val_loss: 0.0576 - val_acc: 0.9822
Epoch 17/20
60000/60000 [=====] - 8s 126us/step - loss: 0.0614 - acc: 0.9803 -
val_loss: 0.0547 - val_acc: 0.9841
Epoch 18/20
60000/60000 [=====] - 8s 131us/step - loss: 0.0599 - acc: 0.9812 -
val_loss: 0.0537 - val_acc: 0.9837
Epoch 19/20
60000/60000 [=====] - 8s 129us/step - loss: 0.0584 - acc: 0.9813 -
val_loss: 0.0577 - val_acc: 0.9816
Epoch 20/20
60000/60000 [=====] - 8s 126us/step - loss: 0.0557 - acc: 0.9818 -
val_loss: 0.0562 - val_acc: 0.9831

```

In [19]:

```

score = model_relu.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])


fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1, nb_epoch+1))

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.05622168810109433  
Test accuracy: 0.9831



In [20]:

```

w_after = model_relu.get_weights()

```

```

w_after = model_relu.get_weights(),


h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()

```



## MLP + RELU activation + adam optimizer + 3 hidden layer

In [19]:

```

# Multilayer perceptron

model_relu = Sequential()
model_relu.add(Dense(261, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.062, seed=None)))
model_relu.add(Dense(87, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.062, seed=None)))
model_relu.add(Dense(29, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.062, seed=None)))
model_relu.add(Dense(output_dim, activation='softmax'))

model_relu.summary()

model_relu.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_relu.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))

```

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/tensorflow/python/framework/op\_def\_library.py:263: colocate\_with (from tensorflow.python.framework.ops) is deprecated and will be removed in a future version.  
Instructions for updating:  
Colocations handled automatically by placer.

| Layer (type)              | Output Shape | Param # |
|---------------------------|--------------|---------|
| dense_1 (Dense)           | (None, 261)  | 204885  |
| dense_2 (Dense)           | (None, 87)   | 22794   |
| dense_3 (Dense)           | (None, 29)   | 2552    |
| dense_4 (Dense)           | (None, 10)   | 300     |
| Total params: 230,531     |              |         |
| Trainable params: 230,531 |              |         |
| Non-trainable params: 0   |              |         |

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/tensorflow/python/ops/math\_ops.py:3066: to\_int32 (from tensorflow.python.ops.math\_ops) is deprecated and will be removed in a future version.  
Instructions for updating:

Use tf.cast instead.  
Train on 60000 samples, validate on 10000 samples  
Epoch 1/20  
60000/60000 [=====] - 5s 81us/step - loss: 0.3684 - acc: 0.8915 - val\_loss: 0.1574 - val\_acc: 0.9520  
Epoch 2/20  
60000/60000 [=====] - 4s 70us/step - loss: 0.1270 - acc: 0.9626 - val\_loss: 0.1077 - val\_acc: 0.9687  
Epoch 3/20  
60000/60000 [=====] - 4s 68us/step - loss: 0.0845 - acc: 0.9746 - val\_loss: 0.0858 - val\_acc: 0.9741  
Epoch 4/20  
60000/60000 [=====] - 4s 67us/step - loss: 0.0634 - acc: 0.9807 - val\_loss: 0.0911 - val\_acc: 0.9706  
Epoch 5/20  
60000/60000 [=====] - 4s 70us/step - loss: 0.0477 - acc: 0.9850 - val\_loss: 0.0780 - val\_acc: 0.9756  
Epoch 6/20  
60000/60000 [=====] - 4s 71us/step - loss: 0.0391 - acc: 0.9873 - val\_loss: 0.0784 - val\_acc: 0.9757  
Epoch 7/20  
60000/60000 [=====] - 4s 70us/step - loss: 0.0301 - acc: 0.9903 - val\_loss: 0.0833 - val\_acc: 0.9762  
Epoch 8/20  
60000/60000 [=====] - 4s 70us/step - loss: 0.0234 - acc: 0.9930 - val\_loss: 0.0800 - val\_acc: 0.9785  
Epoch 9/20  
60000/60000 [=====] - 4s 67us/step - loss: 0.0214 - acc: 0.9931 - val\_loss: 0.0815 - val\_acc: 0.9785  
Epoch 10/20  
60000/60000 [=====] - 4s 70us/step - loss: 0.0182 - acc: 0.9938 - val\_loss: 0.0793 - val\_acc: 0.9785  
Epoch 11/20  
60000/60000 [=====] - 4s 67us/step - loss: 0.0133 - acc: 0.9955 - val\_loss: 0.0819 - val\_acc: 0.9791  
Epoch 12/20  
60000/60000 [=====] - 4s 69us/step - loss: 0.0131 - acc: 0.9958 - val\_loss: 0.0942 - val\_acc: 0.9766  
Epoch 13/20  
60000/60000 [=====] - 4s 68us/step - loss: 0.0129 - acc: 0.9958 - val\_loss: 0.0871 - val\_acc: 0.9776  
Epoch 14/20  
60000/60000 [=====] - 4s 69us/step - loss: 0.0119 - acc: 0.9961 - val\_loss: 0.0891 - val\_acc: 0.9797  
Epoch 15/20  
60000/60000 [=====] - 4s 68us/step - loss: 0.0126 - acc: 0.9961 - val\_loss: 0.0790 - val\_acc: 0.9809  
Epoch 16/20  
60000/60000 [=====] - 4s 69us/step - loss: 0.0085 - acc: 0.9973 - val\_loss: 0.0928 - val\_acc: 0.9806  
Epoch 17/20  
60000/60000 [=====] - 4s 69us/step - loss: 0.0100 - acc: 0.9969 - val\_loss: 0.0914 - val\_acc: 0.9778

```
Epoch 18/20
60000/60000 [=====] - 4s 70us/step - loss: 0.0088 - acc: 0.9971 -
val_loss: 0.0971 - val_acc: 0.9783
Epoch 19/20
60000/60000 [=====] - 4s 68us/step - loss: 0.0053 - acc: 0.9986 -
val_loss: 0.1223 - val_acc: 0.9771
Epoch 20/20
60000/60000 [=====] - 4s 68us/step - loss: 0.0111 - acc: 0.9962 -
val_loss: 0.1049 - val_acc: 0.9796
```

In [20]:


```
score = model_relu.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1, nb_epoch+1))

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

```
Test score: 0.10485377672798267
Test accuracy: 0.9796
```



In [21]:

```
w_after = model_relu.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
h3_w = w_after[4].flatten().reshape(-1,1)
out_w = w_after[6].flatten().reshape(-1,1)



fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 4, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 4, 2)
```

```
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 4, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h3_w,color='y')
plt.xlabel('Hidden Layer 3 ')
plt.show()

plt.subplot(1, 4, 4)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='b')
plt.xlabel('Output Layer ')
plt.show()
```



**MLP + RELU activation + adam optimizer + 3 hidden layer + Batch normalization and Drop out**

In [27]:

```
# Multilayer perceptron

model_relu = Sequential()

model_relu.add(Dense(261, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.062, seed=None)))
model_relu.add(BatchNormalization())
model_relu.add(Dropout(0.5))

model_relu.add(Dense(87, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.062, seed=None)))
model_relu.add(BatchNormalization())
model_relu.add(Dropout(0.5))

model_relu.add(Dense(29, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.062, seed=None)))
model_relu.add(BatchNormalization())
model_relu.add(Dropout(0.5))

model_relu.add(Dense(output_dim, activation='softmax'))

model_relu.summary()

model_relu.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_relu.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

| Layer (type)                                | Output Shape | Param # |
|---|--------------|---------|
| =====                                       | =====        | =====   |
| dense_14 (Dense)                            | (None, 261)  | 204885  |
| batch_normalization_3 (Batch Normalization) | (None, 261)  | 1044    |
| dropout_3 (Dropout)                         | (None, 261)  | 0       |
| dense_15 (Dense)                            | (None, 87)   | 22794   |
| batch_normalization_4 (Batch Normalization) | (None, 87)   | 348     |
| dropout_4 (Dropout)                         | (None, 87)   | 0       |
| dense_16 (Dense)                            | (None, 29)   | 2552    |
| batch_normalization_5 (Batch Normalization) | (None, 29)   | 116     |
| dropout_5 (Dropout)                         | (None, 29)   | 0       |
| dense_17 (Dense)                            | (None, 10)   | 300     |
| =====                                       | =====        | =====   |
| Total params: 232,039                       |              |         |
| Trainable params: 231,285                   |              |         |
| Non-trainable params: 754                   |              |         |

```
Train on 60000 samples, validate on 10000 samples
Epoch 1/20
60000/60000 [=====] - 6s 106us/step - loss: 0.8934 - acc: 0.7249 -
val_loss: 0.2338 - val_acc: 0.9335
Epoch 2/20
60000/60000 [=====] - 5s 86us/step - loss: 0.4027 - acc: 0.8902 -
val_loss: 0.1656 - val_acc: 0.9506
Epoch 3/20
60000/60000 [=====] - 5s 85us/step - loss: 0.3096 - acc: 0.9173 -
val_loss: 0.1398 - val_acc: 0.9585
Epoch 4/20
60000/60000 [=====] - 5s 86us/step - loss: 0.2681 - acc: 0.9302 -
val_loss: 0.1143 - val_acc: 0.9678
Epoch 5/20
60000/60000 [=====] - 5s 84us/step - loss: 0.2352 - acc: 0.9380 -
val_loss: 0.1081 - val_acc: 0.9701
Epoch 6/20
60000/60000 [=====] - 5s 86us/step - loss: 0.2088 - acc: 0.9454
```

```

60000/60000 [=====] - 5s 80us/step - loss: 0.2099 - acc: 0.9454 -
val_loss: 0.0999 - val_acc: 0.9716
Epoch 7/20
60000/60000 [=====] - 5s 86us/step - loss: 0.1978 - acc: 0.9490 -
val_loss: 0.0955 - val_acc: 0.9718
Epoch 8/20
60000/60000 [=====] - 5s 84us/step - loss: 0.1826 - acc: 0.9532 -
val_loss: 0.0936 - val_acc: 0.9734
Epoch 9/20
60000/60000 [=====] - 5s 85us/step - loss: 0.1798 - acc: 0.9538 -
val_loss: 0.0884 - val_acc: 0.9758
Epoch 10/20
60000/60000 [=====] - 5s 85us/step - loss: 0.1686 - acc: 0.9557 -
val_loss: 0.0882 - val_acc: 0.9771
Epoch 11/20
60000/60000 [=====] - 5s 86us/step - loss: 0.1595 - acc: 0.9583 -
val_loss: 0.0822 - val_acc: 0.9770
Epoch 12/20
60000/60000 [=====] - 5s 84us/step - loss: 0.1523 - acc: 0.9610 -
val_loss: 0.0862 - val_acc: 0.9770
Epoch 13/20
60000/60000 [=====] - 5s 83us/step - loss: 0.1489 - acc: 0.9623 -
val_loss: 0.0811 - val_acc: 0.9782
Epoch 14/20
60000/60000 [=====] - 5s 83us/step - loss: 0.1390 - acc: 0.9642 -
val_loss: 0.0761 - val_acc: 0.9790
Epoch 15/20
60000/60000 [=====] - 5s 85us/step - loss: 0.1357 - acc: 0.9656 -
val_loss: 0.0768 - val_acc: 0.9792
Epoch 16/20
60000/60000 [=====] - 5s 84us/step - loss: 0.1304 - acc: 0.9655 -
val_loss: 0.0812 - val_acc: 0.9797
Epoch 17/20
60000/60000 [=====] - 5s 89us/step - loss: 0.1273 - acc: 0.9675 -
val_loss: 0.0804 - val_acc: 0.9790
Epoch 18/20
60000/60000 [=====] - 5s 90us/step - loss: 0.1242 - acc: 0.9682 -
val_loss: 0.0750 - val_acc: 0.9789
Epoch 19/20
60000/60000 [=====] - 5s 86us/step - loss: 0.1230 - acc: 0.9686 -
val_loss: 0.0763 - val_acc: 0.9791
Epoch 20/20
60000/60000 [=====] - 5s 86us/step - loss: 0.1196 - acc: 0.9686 -
val_loss: 0.0754 - val_acc: 0.9814

```

In [28]:

```

score = model_relu.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')


# list of epoch numbers
x = list(range(1,nb_epoch+1))

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

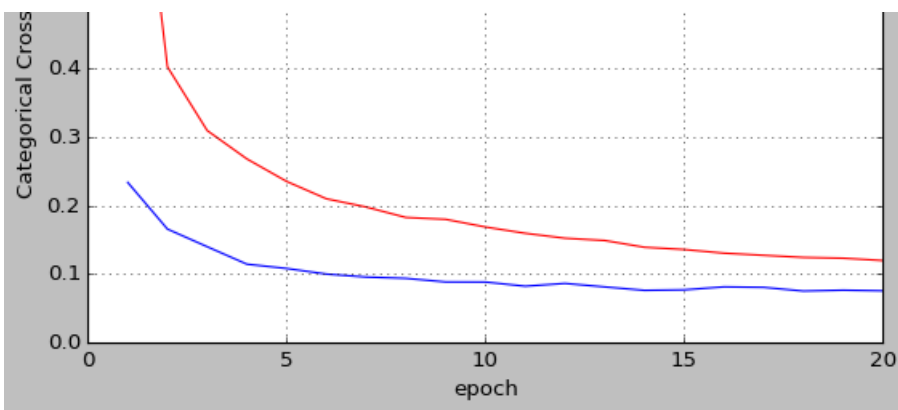
```

Test score: 0.07536343550827004

Test accuracy: 0.9814







In [30]:

```
w_after = model_relu.get_weights()

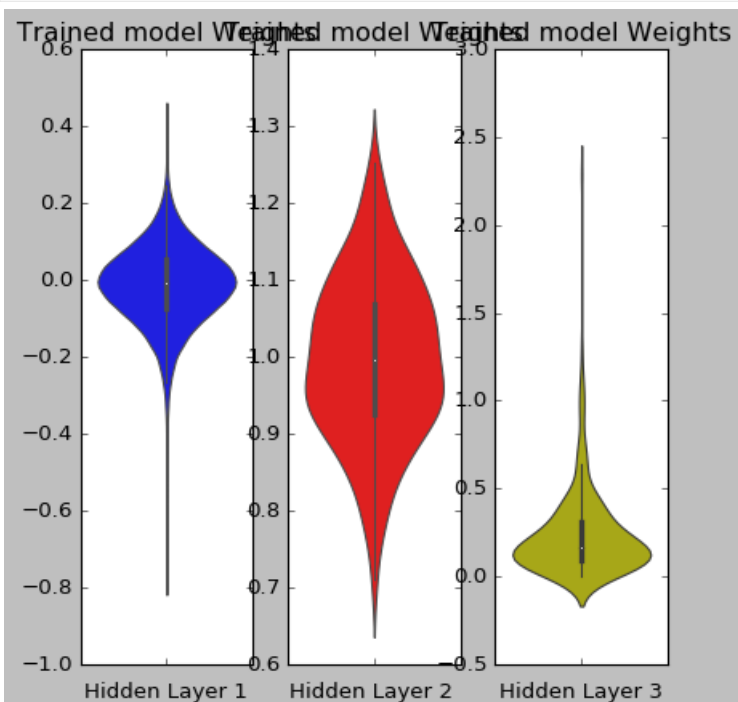
h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
h3_w = w_after[4].flatten().reshape(-1,1)
out_w = w_after[6].flatten().reshape(-1,1)


fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 4, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 4, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 4, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h3_w,color='y')
plt.xlabel('Hidden Layer 3')
plt.show()

plt.subplot(1, 4, 4)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='b')
plt.xlabel('Output Layer ')
plt.show()
```





### MLP + RELU activation + adam optimizer + 5 hidden layer

In [22]:

```
# Multilayer perceptron

model_relu = Sequential()
model_relu.add(Dense(684, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.062, seed=None)))
model_relu.add(Dense(426, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.062, seed=None)))
model_relu.add(Dense(256, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.062, seed=None)))
model_relu.add(Dense(128, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.062, seed=None)))
model_relu.add(Dense(64, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.062, seed=None)))
model_relu.add(Dense(output_dim, activation='softmax'))

model_relu.summary()

model_relu.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_relu.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

| Layer (type)              | Output Shape | Param # |
|---------------------------|--------------|---------|
| dense_5 (Dense)           | (None, 684)  | 536940  |
| dense_6 (Dense)           | (None, 426)  | 291810  |
| dense_7 (Dense)           | (None, 256)  | 109312  |
| dense_8 (Dense)           | (None, 128)  | 32896   |
| dense_9 (Dense)           | (None, 64)   | 8256    |
| dense_10 (Dense)          | (None, 10)   | 650     |
| Total params: 979,864     |              |         |
| Trainable params: 979,864 |              |         |
| Non-trainable params: 0   |              |         |

```

Train on 60000 samples, validate on 10000 samples
Epoch 1/20
60000/60000 [=====] - 15s 249us/step - loss: 0.2423 - acc: 0.9259 - val_loss: 0.1156 - val_acc: 0.9628
Epoch 2/20
60000/60000 [=====] - 14s 238us/step - loss: 0.0878 - acc: 0.9731 - val_loss: 0.0829 - val_acc: 0.9745
Epoch 3/20
60000/60000 [=====] - 14s 238us/step - loss: 0.0602 - acc: 0.9810 - val_loss: 0.0819 - val_acc: 0.9748
Epoch 4/20
60000/60000 [=====] - 14s 236us/step - loss: 0.0451 - acc: 0.9857 - val_loss: 0.0693 - val_acc: 0.9806
Epoch 5/20
60000/60000 [=====] - 15s 242us/step - loss: 0.0342 - acc: 0.9892 - val_loss: 0.0792 - val_acc: 0.9771
Epoch 6/20
60000/60000 [=====] - 14s 242us/step - loss: 0.0299 - acc: 0.9908 - val_loss: 0.0927 - val_acc: 0.9740
Epoch 7/20
60000/60000 [=====] - 14s 240us/step - loss: 0.0253 - acc: 0.9922 - val_loss: 0.0723 - val_acc: 0.9800
Epoch 8/20
60000/60000 [=====] - 15s 243us/step - loss: 0.0231 - acc: 0.9930 - val_loss: 0.0807 - val_acc: 0.9807
Epoch 9/20
60000/60000 [=====] - 15s 242us/step - loss: 0.0218 - acc: 0.9929 - val_loss: 0.0833 - val_acc: 0.9799
Epoch 10/20
60000/60000 [=====] - 14s 241us/step - loss: 0.0184 - acc: 0.9942 - val_loss: 0.1071 - val_acc: 0.9752
Epoch 11/20
60000/60000 [=====] - 14s 240us/step - loss: 0.0180 - acc: 0.9945 - val_loss: 0.0822 - val_acc: 0.9793
Epoch 12/20
60000/60000 [=====] - 14s 238us/step - loss: 0.0132 - acc: 0.9960 - val_loss: 0.0866 - val_acc: 0.9809
Epoch 13/20
60000/60000 [=====] - 14s 239us/step - loss: 0.0151 - acc: 0.9955 - val_loss: 0.0919 - val_acc: 0.9790
Epoch 14/20
60000/60000 [=====] - 14s 237us/step - loss: 0.0120 - acc: 0.9966 - val_loss: 0.1011 - val_acc: 0.9787
Epoch 15/20
60000/60000 [=====] - 14s 241us/step - loss: 0.0150 - acc: 0.9954 - val_loss: 0.1007 - val_acc: 0.9776
Epoch 16/20
60000/60000 [=====] - 14s 238us/step - loss: 0.0135 - acc: 0.9957 - val_loss: 0.0896 - val_acc: 0.9826
Epoch 17/20
60000/60000 [=====] - 14s 240us/step - loss: 0.0124 - acc: 0.9968 - val_loss: 0.0818 - val_acc: 0.9825
Epoch 18/20
60000/60000 [=====] - 14s 241us/step - loss: 0.0111 - acc: 0.9970 - val_loss: 0.0772 - val_acc: 0.9818
Epoch 19/20
60000/60000 [=====] - 14s 241us/step - loss: 0.0105 - acc: 0.9971 - val_loss: 0.0911 - val_acc: 0.9799
Epoch 20/20
60000/60000 [=====] - 14s 241us/step - loss: 0.0103 - acc: 0.9971 - val_loss: 0.0923 - val_acc: 0.9813

```

In [23]:

```

score = model_relu.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))


vy = history.history['val_loss']
ty = history.history['loss']

```

```
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.09227842169566275

Test accuracy: 0.9813



In [26]:

```
w_after = model_relu.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
h3_w = w_after[4].flatten().reshape(-1,1)
h4_w = w_after[6].flatten().reshape(-1,1)
out_w = w_after[8].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 6, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 6, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')





plt.subplot(1, 6, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h3_w,color='y')
plt.xlabel('Hidden Layer 3')
plt.show()

plt.subplot(2, 6, 4)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h3_w,color='b')
plt.xlabel('Hidden Layer 4')
plt.show()

plt.subplot(2, 6, 5)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h3_w,color='r')
plt.xlabel('Hidden Layer 5')
plt.show()

plt.subplot(2, 6, 6)
plt.title("Trained model Weights")
```

```
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



## MLP + RELU activation + adam optimizer + 5 hidden layer + Batch normalization and Drop out

In [27]:

```
# Multilayer perceptron

model_relu = Sequential()
model_relu.add(Dense(684, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.062, seed=None)))
model_relu.add(BatchNormalization())
model_relu.add(Dropout(0.5))

model_relu.add(Dense(426, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.062, seed=None)))
model_relu.add(BatchNormalization())
model_relu.add(Dropout(0.5))

model_relu.add(Dense(256, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.062, seed=None)))
model_relu.add(BatchNormalization())
model_relu.add(Dropout(0.5))

model_relu.add(Dense(128, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.062, seed=None)))
model_relu.add(BatchNormalization())
model_relu.add(Dropout(0.5))

model_relu.add(Dense(64, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.062, seed=None)))
model_relu.add(BatchNormalization())
model_relu.add(Dropout(0.5))

model_relu.add(Dense(output_dim, activation='softmax'))

model_relu.summary()

model_relu.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_relu.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow\_backend.py:3445: calling dropout (from tensorflow.python.ops.nn\_ops) with keep\_prob is deprecated and will be removed in a future version.

Instructions for updating:

Please use `rate` instead of `keep\_prob`. Rate should be set to `rate = 1 - keep\_prob`.

| Layer (type)                                | Output Shape | Param # |
|---|--------------|---------|
| dense_11 (Dense)                            | (None, 684)  | 536940  |
| batch_normalization_1 (Batch Normalization) | (None, 684)  | 2736    |
| dropout_1 (Dropout)                         | (None, 684)  | 0       |
| dense_12 (Dense)                            | (None, 426)  | 291810  |
| batch_normalization_2 (Batch Normalization) | (None, 426)  | 1704    |
| dropout_2 (Dropout)                         | (None, 426)  | 0       |
| dense_13 (Dense)                            | (None, 256)  | 109312  |
| batch_normalization_3 (Batch Normalization) | (None, 256)  | 1024    |
| dropout_3 (Dropout)                         | (None, 256)  | 0       |
| dense_14 (Dense)                            | (None, 128)  | 32896   |
| batch_normalization_4 (Batch Normalization) | (None, 128)  | 512     |
| dropout_4 (Dropout)                         | (None, 128)  | 0       |

|   |            |      |
|---|------------|------|
| dense_15 (Dense)                            | (None, 64) | 8256 |
| batch_normalization_5 (Batch Normalization) | (None, 64) | 256  |
| dropout_5 (Dropout)                         | (None, 64) | 0    |
| dense_16 (Dense)                            | (None, 10) | 650  |

=====  
 Total params: 986,096  
 Trainable params: 982,980  
 Non-trainable params: 3,116

Train on 60000 samples, validate on 10000 samples

Epoch 1/20  
 60000/60000 [=====] - 19s 319us/step - loss: 1.0018 - acc: 0.6862 - val\_loss: 0.2397 - val\_acc: 0.9278

Epoch 2/20  
 60000/60000 [=====] - 18s 293us/step - loss: 0.3527 - acc: 0.9032 - val\_loss: 0.1574 - val\_acc: 0.9541

Epoch 3/20  
 60000/60000 [=====] - 18s 296us/step - loss: 0.2587 - acc: 0.9309 - val\_loss: 0.1173 - val\_acc: 0.9672

Epoch 4/20  
 60000/60000 [=====] - 18s 295us/step - loss: 0.2184 - acc: 0.9431 - val\_loss: 0.1031 - val\_acc: 0.9708

Epoch 5/20  
 60000/60000 [=====] - 18s 292us/step - loss: 0.1864 - acc: 0.9505 - val\_loss: 0.0994 - val\_acc: 0.9732

Epoch 6/20  
 60000/60000 [=====] - 18s 292us/step - loss: 0.1689 - acc: 0.9556 - val\_loss: 0.0965 - val\_acc: 0.9748

Epoch 7/20  
 60000/60000 [=====] - 18s 293us/step - loss: 0.1579 - acc: 0.9586 - val\_loss: 0.0828 - val\_acc: 0.9788

Epoch 8/20  
 60000/60000 [=====] - 18s 292us/step - loss: 0.1486 - acc: 0.9619 - val\_loss: 0.0818 - val\_acc: 0.9777

Epoch 9/20  
 60000/60000 [=====] - 18s 295us/step - loss: 0.1373 - acc: 0.9646 - val\_loss: 0.0780 - val\_acc: 0.9785

Epoch 10/20  
 60000/60000 [=====] - 18s 294us/step - loss: 0.1248 - acc: 0.9674 - val\_loss: 0.0773 - val\_acc: 0.9802

Epoch 11/20  
 60000/60000 [=====] - 17s 291us/step - loss: 0.1194 - acc: 0.9688 - val\_loss: 0.0809 - val\_acc: 0.9779

Epoch 12/20  
 60000/60000 [=====] - 18s 292us/step - loss: 0.1192 - acc: 0.9686 - val\_loss: 0.0676 - val\_acc: 0.9816

Epoch 13/20  
 60000/60000 [=====] - 17s 290us/step - loss: 0.1112 - acc: 0.9709 - val\_loss: 0.0748 - val\_acc: 0.9804

Epoch 14/20  
 60000/60000 [=====] - 18s 294us/step - loss: 0.1096 - acc: 0.9704 - val\_loss: 0.0703 - val\_acc: 0.9815

Epoch 15/20  
 60000/60000 [=====] - 18s 296us/step - loss: 0.1012 - acc: 0.9733 - val\_loss: 0.0710 - val\_acc: 0.9821

Epoch 16/20  
 60000/60000 [=====] - 18s 296us/step - loss: 0.0958 - acc: 0.9747 - val\_loss: 0.0759 - val\_acc: 0.9799

Epoch 17/20  
 60000/60000 [=====] - 18s 295us/step - loss: 0.0937 - acc: 0.9760 - val\_loss: 0.0668 - val\_acc: 0.9821

Epoch 18/20  
 60000/60000 [=====] - 18s 294us/step - loss: 0.0882 - acc: 0.9766 - val\_loss: 0.0648 - val\_acc: 0.9834

Epoch 19/20  
 60000/60000 [=====] - 18s 295us/step - loss: 0.0855 - acc: 0.9778 - val\_loss: 0.0694 - val\_acc: 0.9823

Epoch 20/20  
 60000/60000 [=====] - 17s 291us/step - loss: 0.0857 - acc: 0.9774 - val\_loss: 0.0623 - val\_acc: 0.9834

In [20]:

```
score = model_relu.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

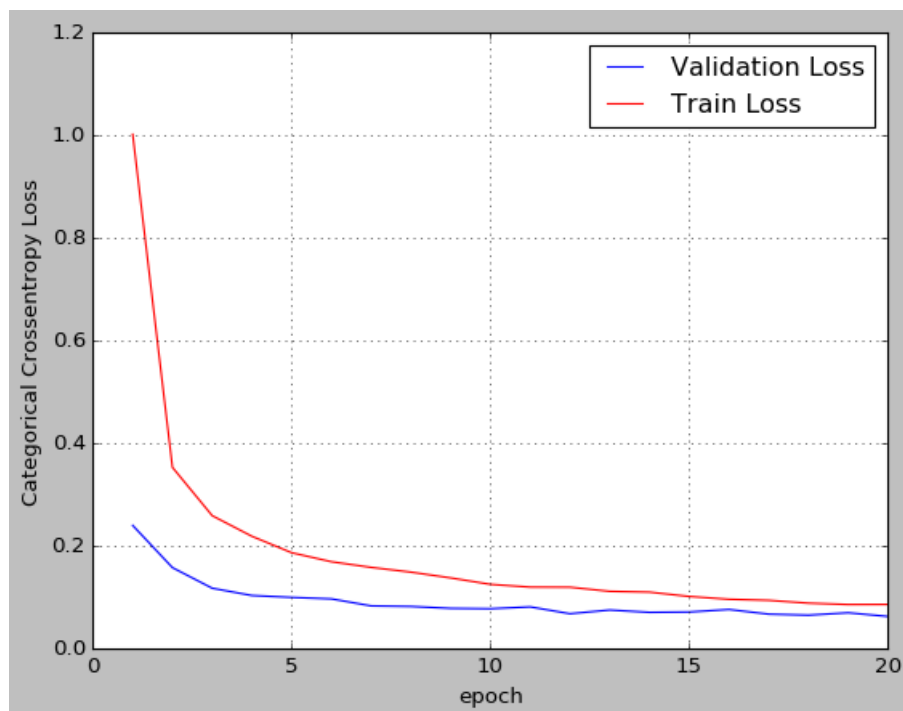
fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1, nb_epoch+1))

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.06228468422386795

Test accuracy: 0.9834



In [29]:

```
w_after = model_relu.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
h3_w = w_after[4].flatten().reshape(-1,1)
h4_w = w_after[6].flatten().reshape(-1,1)
out_w = w_after[8].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 6, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 6, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 6, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h3_w,color='y')
plt.xlabel('Hidden Layer 3')
plt.show()

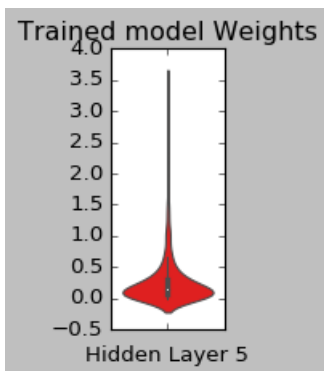
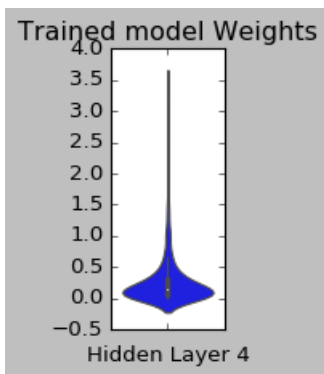

plt.subplot(2, 6, 4)
```



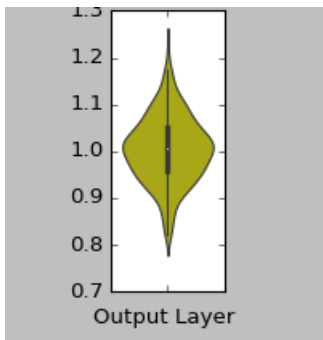
```
plt.subplot(2, 6, 4)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h3_w,color='b')
plt.xlabel('Hidden Layer 4')
plt.show()

plt.subplot(2, 6, 5)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h3_w,color='r')
plt.xlabel('Hidden Layer 5')
plt.show()

plt.subplot(2, 6, 6)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



Trained model Weights



In [7]:

```
#code copied from -http://zetcode.com/python/prettytable/
from prettytable import PrettyTable
x = PrettyTable()
x.field_names = ["No. of Hidden layers", "Normalization and dropout applied?", "Accuracy"]
x.add_row([2, "No", 0.9818])
x.add_row([2, "Yes", 0.9831])
x.add_row([3, "No", 0.9796])
x.add_row([3, "Yes", 0.9814])
x.add_row([5, "No", 0.9813])
x.add_row([5, "Yes", 0.9834])
print(x)
```

| No. of Hidden layers | Normalization and dropout applied? | Accuracy |
|----------------------|------------------------------------|----------|
| 2                    | No                                 | 0.9818   |
| 2                    | Yes                                | 0.9831   |
| 3                    | No                                 | 0.9796   |
| 3                    | Yes                                | 0.9814   |
| 5                    | No                                 | 0.9813   |
| 5                    | Yes                                | 0.9834   |

**The best MLP architecture for MNIST Data Set is when we applied 5 hidden layers with Normalization and Drop Out applied and accuracy obtained is 98.34 percent**